





# Parallelism & Concurrency

# Goals For Today



- Introduce the Ideas of Parallelism & Concurrency

# Don't Forget!



- HW7 Due Tonight
- MP2 Due 3/4
- MP3 Releases 3/3

# Don't Forget!



- HW7 Due Tonight
- MP2 Due 3/4
- MP3 Releases 3/3

# But First...



Could you explain how lowercase works for characters?



# But First...



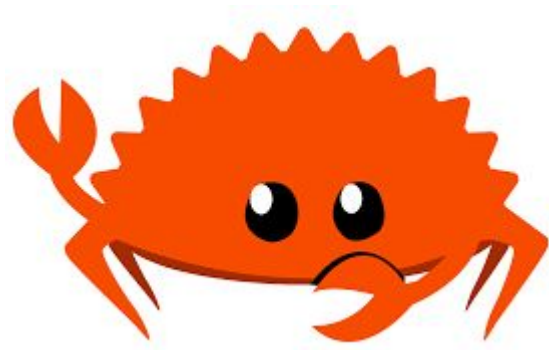
Could you explain how lowercase works for characters?

Converting

- `.to_uppercase()`
- `.to_lowercase()`

Checking

- `.is_uppercase()`
- `.is_lowercase()`



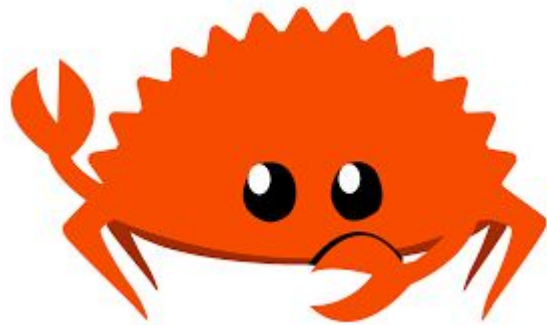
# But First...



Could use another refresher on `.unwrap`, also why we use `.as_ref` instead of `&`?

Unwrap refresher:

- It's very common that certain operations may fail (`Result`)
  - Like if you try to read a file that doesn't exist
- We don't necessarily want to panic (and thus end execution), so instead of returning some type we can return `Result<T>`
- Then, we can check if what we return is `Ok(T)` or `Err(E*)`
- If it is `Ok(T)`, we probably want to use whatever `T` is
- So we can `.unwrap()` it



\*We use `E` to represent an error type



# But First...



Could use another refresher on `.unwrap`, also why we use `.as_ref` instead of `&`?

Unwrap refresher:

- It's very common that certain operations may ~~fail (Result)~~ return nothing (Option)
  - Like if you try to read a file that is empty
- We don't necessarily want to panic (and thus end execution), so instead of returning some type we can return ~~Result~~ <sup>Option</sup> `<T>`
- Then, we can check if what we return is ~~Ok~~ <sup>Some</sup> `(T)` or ~~Err(E\*)~~ <sup>None</sup>
- If it is ~~Ok~~ <sup>Some</sup> `(T)`, we probably want to use whatever `T` is
- So we can `.unwrap()` it



# But First...



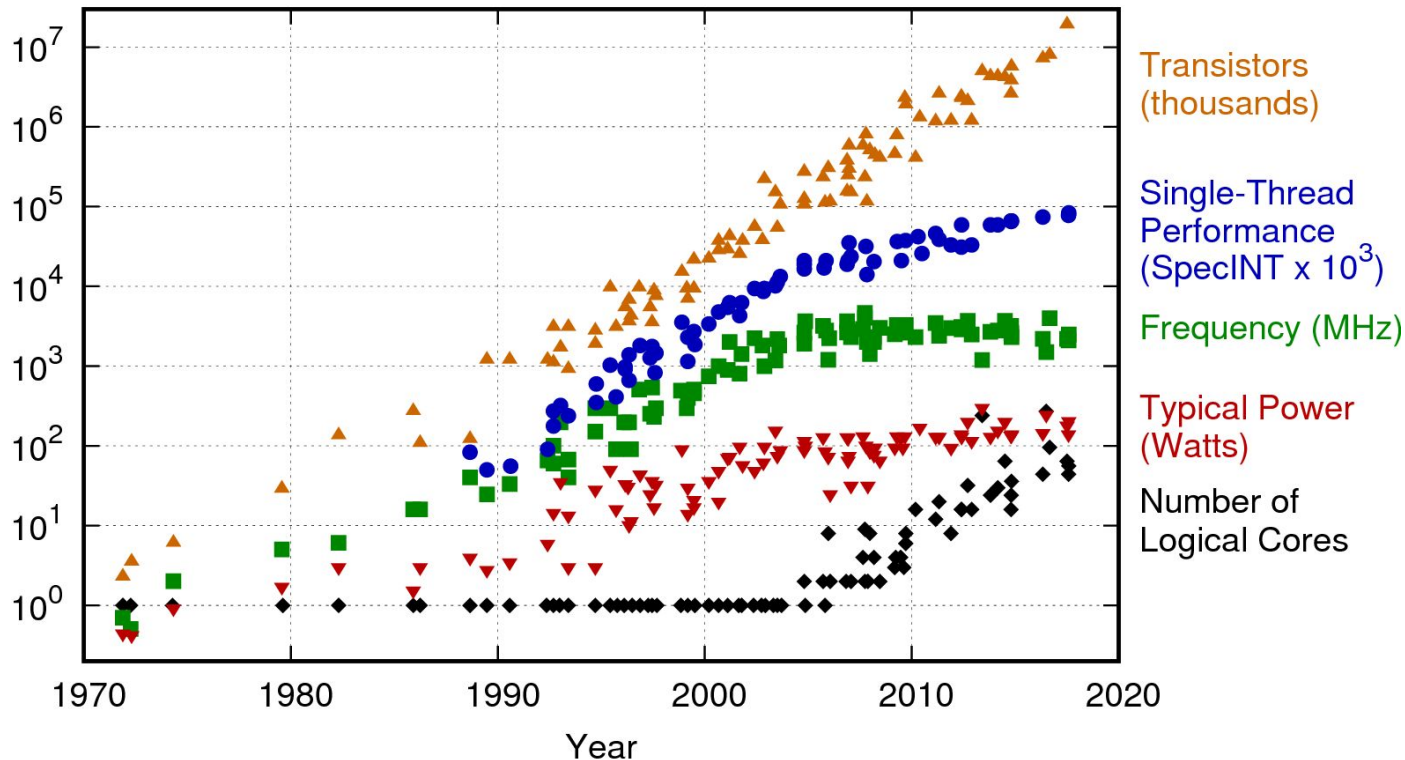
Could you go over how to use clone? How does it differ to trying to use a reference?



# Why Concurrency?



42 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2017 by K. Rupp

# Why Concurrency?



*"Scalable algorithms and libraries can be the best legacy we can leave behind from this era"*

-David Kirk and Wen-mei W. Hwu's ECE 408 Slides

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

read

0

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

Write

1

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

0

Process 2:

```
global_data = 10;
```

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

$x=0$

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

0

read

Process 2:

```
global_data = 10;
```



# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

`x=0`

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

10

Process 2:

```
global_data = 10;
```

write

10

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

$x=0$

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

write

1

Process 2:

```
global_data = 10;
```

# Why Not Concurrency?



Concurrency is **hard**.

Process 1:

`x=0`

```
let x = global_data;  
global_data = x + 1;
```

Global  
Data

write

1

Process 2:

```
global_data = 10;  
// I think it is == 10!
```

# Why Not Concurrency?

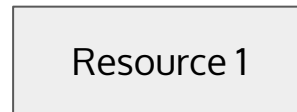
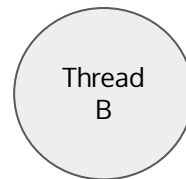
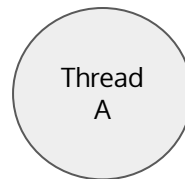


Concurrency is **hard**.

Common issues:

1. **Race Conditions**
2. Deadlocks
3. Livelocks
4. Starvation

Race	Thread A	Thread B	Resources
Time 1	Accesses 1	Accesses 1	Resource 1
Time 2	Resource 1 = 2	Resource 1 = 3	Resource 2
Time 3	Resource 1 += 1	Resource 1 += 1	Resource 3
Time 4	What is Resource 1?	What is Resource 1?	Resource 4



# Why Not Concurrency?

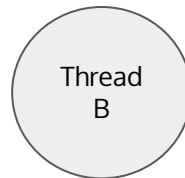
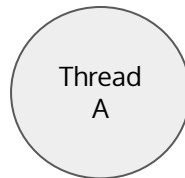


Concurrency is **hard**.

Common issues:

1. Race Conditions
2. **Deadlocks**
3. Livelocks
4. Starvation

Deadlock Example	Thread A	Thread B	Resources
Time 1	Accesses 1 & 2	Accesses 3 & 4	Resource 1
Time 2	Attempts to Access 3	Denies Access to 3	Resource 2
Time 3	Denies Access to 1	Attempts Access to 1	Resource 3
Time 4	Try Again (Deadlock)	Try Again (Deadlock)	Resource 4



# Why Not Concurrency?

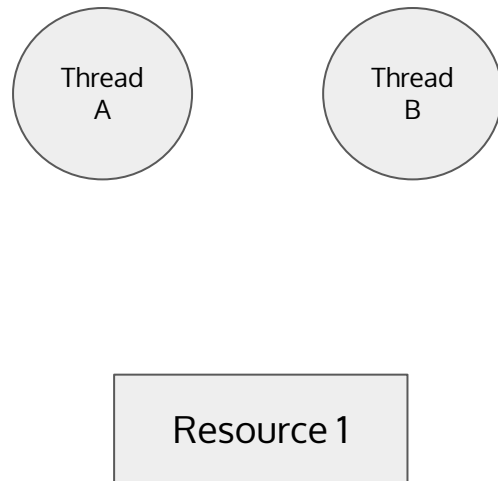


Concurrency is **hard**.

Common issues:

1. Race Conditions
2. Deadlocks
3. **Livelocks**
4. Starvation

Livelock	Thread A	Thread B	Resources
Time 1	Accesses Resource 1	Requests Resource 1	Resource 1
Time 2	Requests Resource 1	Accesses Resource 1	Resource 2
Time 3	Accesses Resource 1	Requests Resource 1	Resource 3
Time 4	And so on.. (livelock)	And so on.. (livelock)	Resource 4



# Why Not Concurrency?

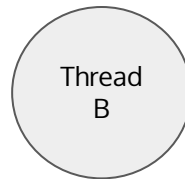
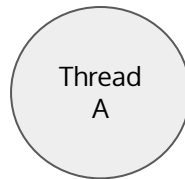


Concurrency is **hard**.

Common issues:

1. Race Conditions
2. Deadlocks
3. Livelocks
4. **Starvation**

Starvation	Thread A	Thread B	Resources
Time 1	Accesses 1, 2, 3, & 4	Politely Yields to A	Resource 1
Time 2	Uses 1, 2, 3, & 4	Requests Access	Resource 2
Time 3	Uses 1, 2, 3, & 4	Requests Access	Resource 3
Time 4	Continued Use (Starvation)	Continued Wait (Starvation)	Resource 4



# Why Not Concurrency?



Concurrency is **hard**.



# Concurrency!



Concurrency is **hard**.

But Rust makes it **Easy**.

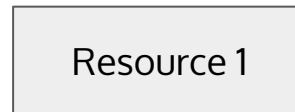
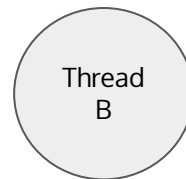
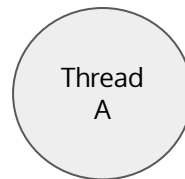
# Concurrency!



Concurrency is hard.

But Rust makes it Easy.

Race	Thread A	Thread B	Resources
Time 1	Accesses 1	Accesses 1	Resource 1
Time 2	Resource 1 = 2	Resource 1 = 3	Resource 2
Time 3	Resource 1 += 1	Resource 1 += 1	Resource 3
Time 4	What is Resource 1?	What is Resource 1?	Resource 4



# Concurrency!

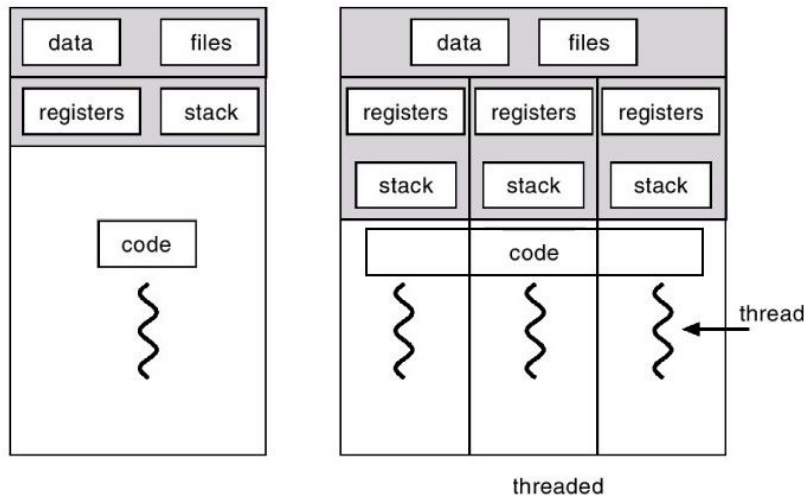
Concurrency is hard.

But Rust makes it Easy.

An **operating system** manages many **processes**\*

**Processes** may have many **threads**

**Threads** are the execution of code.



\*There's some distinction between processes and programs. Generally, we refer to programs as the code while processes are the things that execute that code. That is, your OS creates a process to execute a program.

## Concurrency vs Parallelism

- **Concurrency** is being able to switch tasks when we have some downtime
- **Parallelism** is being able to execute multiple tasks at the same time

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel Optane memory access	<10 $\mu$ s	7 hrs
NVMe SSD I/O	~25 $\mu$ s	17 hrs
SSD I/O	50–150 $\mu$ s	1.5–4 days
Rotational disk I/O	1–10 ms	1–9 months
Internet call: SF to NYC	65 ms	5 years
Internet call: SF to Hong Kong	141 ms	11 years

# Concurrency!



Effective concurrency is highly dependent on effective communication.

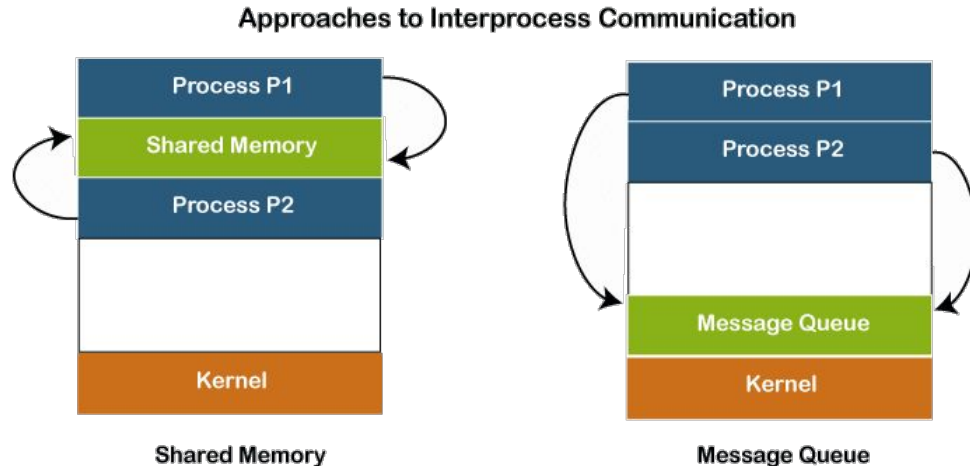
There's two major ways of communicating that we'll explore.

# Concurrency!

Effective concurrency is highly dependent on effective communication.

There's two major ways of communicating that we'll explore.

## Message Passing and Shared Memory



# Concurrency Examples



Ok Eustis very cool but how do we actually do it?

# Concurrency Examples



```
use std::thread;
use std::time::Duration;

fn main() {
    let handle = thread::spawn(|| {
        for i in 1..10 {
            println!("hi number {} from the spawned thread!", i);
            thread::sleep(Duration::from_millis(1));
        }
    });

    for i in 1..5 {
        println!("hi number {} from the main thread!", i);
        thread::sleep(Duration::from_millis(1));
    }
}
```



# Concurrency Examples - Closures



```
use std::thread;
use std::time::Duration;

fn main() {
    let closure = |num| -> i32 {
        println!("calculating slowly...");
        thread::sleep(Duration::from_secs(2));
        return num + 1
    };

    println!("{}", closure(3));
}
```

Closures: **Anonymous** Functions that Can Capture Their Environment

# Concurrency Examples - Closures



```
use std::thread;
use std::time::Duration;

fn main() {
    let closure = |num| -> i32 {
        println!("calculating slowly...");
        thread::sleep(Duration::from_secs(2));
        return num + 1
    };

    println!("{}", closure(3));
}
```

Closures: Anonymous Functions that Can **Capture Their Environment**

# Concurrency Examples - Closures



```
use std::thread;
use std::time::Duration;

fn main() {
    let closure = |num| -> i32 {
        println!("calculating slowly...");
        thread::sleep(Duration::from_secs(2));
        return num + 1
    };

    println!("{}", closure(3));
}
```

Closures: Anonymous Functions that Can **Capture Their Environment**

- closures can capture values from the scope in which they're defined

# Concurrency Examples - Closures



```
use std::thread;
use std::time::Duration;

fn main() {
    let closure = |num| -> i32 {
        println!("calculating slowly...");
        thread::sleep(Duration::from_secs(2));
        return num + 1
    };

    println!("{}", closure(3));
}
```

# Concurrency Examples - Closures



```
use std::thread;
use std::time::Duration;

fn main() {
    let closure = |num| -> i32 {
        println!("calculating slowly...");
        thread::sleep(Duration::from_secs(2));
        return num + 1
    };

    println!("{}", closure(3));
}
```



```
fn add_one_v1 (x: u32) -> u32 { x + 1 }
let add_one_v2 = |x: u32| -> u32 { x + 1 };
let add_one_v3 = |x|           { x + 1 };
let add_one_v4 = |x|           x + 1 ;
```

# Concurrency Examples



```
use std::thread;
use std::time::Duration;

fn main() {
    let handle = thread::spawn(|| {
        for i in 1..10 {
            println!("hi number {} from the spawned thread!", i);
            thread::sleep(Duration::from_millis(1));
        }
    });

    for i in 1..5 {
        println!("hi number {} from the main thread!", i);
        thread::sleep(Duration::from_millis(1));
    }
}
```

```
hi number 1 from the main thread!
hi number 1 from the spawned thread!
hi number 2 from the main thread!
hi number 2 from the spawned thread!
hi number 3 from the main thread!
hi number 3 from the spawned thread!
hi number 4 from the main thread!
hi number 4 from the spawned thread!
```

```
hi number 1 from the main thread!
hi number 1 from the spawned thread!
hi number 2 from the spawned thread!
hi number 3 from the spawned thread!
hi number 4 from the spawned thread!
hi number 5 from the spawned thread!
hi number 2 from the main thread!
hi number 6 from the spawned thread!
hi number 3 from the main thread!
hi number 7 from the spawned thread!
hi number 4 from the main thread!
hi number 8 from the spawned thread!
```

# Concurrency Examples



```
use std::thread;
use std::time::Duration;

fn main() {
    let handle = thread::spawn(|| {
        for i in 1..10 {
            println!("hi number {} from the spawned thread!", i);
            thread::sleep(Duration::from_millis(1));
        }
    });

    for i in 1..5 {
        println!("hi number {} from the main thread!", i);
        thread::sleep(Duration::from_millis(1));
    }

    handle.join().unwrap()
}
```

# Concurrency Examples



```
use std::thread;
use std::time::Duration;

fn main() {
    let handle = thread::spawn(|| {
        for i in 1..9 {
            println!("hi number {} from the spawned thread!", i);
            thread::sleep(Duration::from_millis(1));
        }
    });

    handle.join().unwrap();

    for i in 1..5 {
        println!("hi number {} from the main thread!", i);
        thread::sleep(Duration::from_millis(1));
    }
}
```



# That's All Folks!



See you on Thursday :)

# But First...



Could use another refresher on `.unwrap`, also why we use `.as_ref` instead of `&`?

Why do we use `.as_ref` instead of `&`?

- This one is a sneaky difference that enters a large tangent, but here we go.
- When we are talking about the difference, we really care about the difference in the behavior of the **Borrow** and **AsRef** traits.



\*The Docs: <https://doc.rust-lang.org/std/convert/trait.AsRef.html>

# But First...

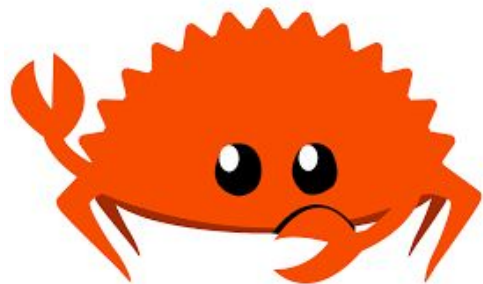


Could use another refresher on `.unwrap`, also why we use `.as_ref` instead of `&`?

Why do we use `.as_ref` instead of `&`?

From the docs\*:

- **AsRef** has the same signature as **Borrow**, but **Borrow** is different in few aspects:
  - Unlike `AsRef`, `Borrow` has a blanket impl for any `T`, and can be used to accept either a reference or a value.
  - `Borrow` also requires that `Hash`, `Eq` and `Ord` for borrowed value are equivalent to those of the owned value. For this reason, if you want to borrow only a single field of a struct you can implement `AsRef`, but not `Borrow`.
    - `(x == y) == (&x == &y)`



\*The Docs: <https://doc.rust-lang.org/std/convert/trait.AsRef.html>

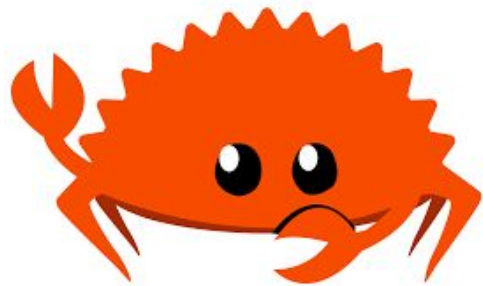
# But First...



Could use another refresher on .unwrap, also why we use .as\_ref instead of &?

Why do we use .as\_ref instead of &?

```
fn is_hello<T: AsRef<str>>(s: T) {  
    assert_eq!("hello", s.as_ref());  
}  
  
let s = "hello";  
is_hello(s);  
  
let s = "hello".to_string();  
is_hello(s);
```



# But First...



Could use another refresher on `.unwrap`, also why we use `.as_ref` instead of `&`?

Why do we use `.as_ref` instead of `&`?

Choose **Borrow** when you want to abstract over different kinds of borrowing, or when you're building a data structure that treats owned and borrowed values in equivalent ways, such as hashing and comparison.

Choose **AsRef** when you want to convert something to a reference directly, and you're writing generic code.

